


Please check your email address for a confirmation link

Resend confirmation email

FLUX.1 Kontext [dev]

Image editing and manipulation model guidance-distilled from FLUX.1 Kontext [pro], [\[blog\]](#) [\[model\]](#)



Reuse this image

Please check your email address for a confirmation link [Resend confirmation email](#)

Stable Diffusion 2.1 Demo



Stable Diffusion 2.1 is the latest text-to-image model from StabilityAI. [Access Stable Diffusion 1 Space here](#)

For faster generation and API access you can try [DreamStudio Beta](#).

girl sitting in a chair

Generate image

Enter a negative prompt



Demo

Past generations


Illusion Diffusion HQ

Generate stunning high quality illusion artwork with Stable Diffusion

Illusion Diffusion is back up with a safety checker! Because I have been asked, if you would like to support me, consider using deforum.studio

A space by AP [Follow me on Twitter](#) with big contributions from [multimodalart](#)

This project works by using [Monster Labs QR Control Net](#). Given a prompt and your pattern, we use a QR code conditioned controlnet to create a stunning illusion! Credit to: [MrUgleh](#) for discovering the workflow :)



!nvidia-smi



Tue Jul 29 17:21:48 2025

NVIDIA-SMI 550.54.15				Driver Version	
GPU	Name		Persistence-M	Bus-Id	
Fan	Temp	Perf	Pwr:Usage/Cap		
0	Tesla	T4	Off	0000	
N/A	54C	P8	10W / 70W		

Processes:						
GPU	GI	CI	PID	Type	Process name	
	ID	ID				
=====						
No running processes found						

beautifull girl.png X



1.Text Generation Model(Model-GPT2)

```
from transformers import pipeline
```

```
generator = pipeline("text-generation", model= 'gpt2')
response = generator('What is Hugging Face', max_length
print(response[0]['generated_text'])
```



```
/usr/local/lib/python3.11/dist-packages/huggingfac
The secret `HF_TOKEN` does not exist in your Colab
To authenticate with the Hugging Face Hub, create
You will be able to reuse this secret in all of yo
Please note that authentication is recommended but
warnings.warn(
```

```
config.json: 100%          665/665 [00:00<00:00, 58.6kB/s]
```

```
model.safetensors: 100%   548M/548M [00:10<00:00, 97.2MB/s]
```

```
generation_config.json: 100%  124/124 [00:00<00:00, 6.36kB/s]
```

```
tokenizer_config.json: 100%   26.0/26.0 [00:00<00:00, 2.95kB/s]
```

```
vocab.json: 100%          1.04M/1.04M [00:00<00:00, 7.01MB/s]
```

```
merges.txt: 100%          456k/456k [00:00<00:00, 5.98MB/s]
```

```
tokenizer.json: 100%       1.36M/1.36M [00:00<00:00, 6.24MB/s]
```

```
Device set to use cuda:0
```

```
Truncation was not explicitly activated but `max_l
Setting `pad_token_id` to `eos_token_id`:50256 for
Both `max_new_tokens` (=256) and `max_length` (=40)
What is Hugging Face?
```

```
You can use the Hugging Face tool to find out what
```

```
How does Hugging Face work?
```

```
The Hugging Face tool is a great tool for finding
```

```
Your friends are also a great tool to know, becaus
```

```
How does the Hugging Face test work?
```

```
The Hugging Face test is for finding out who we ar
```

```
How does the Hugging Face test work in real life?
```

```
You can use the Hugging Face tool to find out what
```

✓ 2.Named Entity Recognition

```
from transformers import pipeline
```

```
ner = pipeline('ner', grouped_entities=True)
entities = ner('Hugging face is based on NYC and partn
print(entities)
```

➞ No model was supplied, defaulted to dbmdz/bert-lar
Using a pipeline without specifying a model name a

```
config.json: 100%          998/998 [00:00<00:00, 106kB/s]

model.safetensors: 100%  1.33G/1.33G [00:25<00:00, 28.1MB/

Some weights of the model checkpoint at dbmdz/bert
- This IS expected if you are initializing BertFor
- This IS NOT expected if you are initializing Ber

tokenizer_config.json: 100%  60.0/60.0 [00:00<00:00, 3.61kB/s]

vocab.txt:      213k/? [00:00<00:00, 8.52MB/s]

Device set to use cuda:0
[{'entity_group': 'LOC', 'score': np.float32(0.995
/usr/local/lib/python3.11/dist-packages/transforme
warnings.warn(
/usr/local/lib/python3.11/dist-packages/torch/nn/m
return forward call(*args, **kwargs)
```

```
from transformers import pipeline
```

```
ner = pipeline('ner', grouped_entities=True)
entities = ner('My name is Priyanka learning Genai')
print(entities)
```

➞ No model was supplied, defaulted to dbmdz/bert-lar
Using a pipeline without specifying a model name a
Some weights of the model checkpoint at dbmdz/bert
- This IS expected if you are initializing BertFor
- This IS NOT expected if you are initializing Ber
Device set to use cuda:0
[{'entity_group': 'PER', 'score': np.float32(0.925

✓ 3.Sentiment Analysis

```
!pip install transformers datasets
```

➞ Requirement already satisfied: transformers in /usr
Requirement already satisfied: datasets in /usr/lo
Requirement already satisfied: filelock in /usr/lo
Requirement already satisfied: huggingface-hub<1.0
Requirement already satisfied: numpy>=1.17 in /usr
Requirement already satisfied: packaging>=20.0 in
Requirement already satisfied: pyyaml>=5.1 in /usr
Requirement already satisfied: regex!=2019.12.17 i
Requirement already satisfied: requests in /usr/lo
Requirement already satisfied: tokenizers<0.22,>=0
Requirement already satisfied: safetensors>=0.4.3
Requirement already satisfied: tqdm>=4.27 in /usr/
Requirement already satisfied: pyarrow>=15.0.0 in
Requirement already satisfied: dill<0.3.9,>=0.3.0
Requirement already satisfied: pandas in /usr/loca

```
Requirement already satisfied: xxhash in /usr/loca
Requirement already satisfied: multiprocessing<0.70.1
Requirement already satisfied: fsspec<=2025.3.0,>=
Requirement already satisfied: aiohttp!=4.0.0a0,!=
Requirement already satisfied: typing-extensions>=
Requirement already satisfied: hf-xet<2.0.0,>=1.1.
Requirement already satisfied: charset-normalizer<
Requirement already satisfied: idna<4,>=2.5 in /us
Requirement already satisfied: urllib3<3,>=1.21.1
Requirement already satisfied: certifi>=2017.4.17
Requirement already satisfied: python-dateutil>=2.
Requirement already satisfied: pytz>=2020.1 in /us
Requirement already satisfied: tzdata>=2022.7 in /
Requirement already satisfied: aiohappyeyeballs>=2
Requirement already satisfied: aiosignal>=1.4.0 in
Requirement already satisfied: attrs>=17.3.0 in /u
Requirement already satisfied: frozenlist>=1.1.1 i
Requirement already satisfied: multidict<7.0,>=4.5
Requirement already satisfied: propcache>=0.2.0 in
Requirement already satisfied: yarll<2.0,>=1.17.0 i
Requirement already satisfied: six>=1.5 in /usr/lo
```

```
from transformers import pipeline
```

```
classifier = pipeline('sentiment-analysis')
result = classifier('I love Hugging Face library very
print(result)
```



```
No model was supplied, defaulted to distilbert/dis
Using a pipeline without specifying a model name a
config.json: 100%          629/629 [00:00<00:00, 52.7kB/s]

model.safetensors: 100% 268M/268M [00:03<00:00, 111MB/s]

tokenizer_config.json: 100% 48.0/48.0 [00:00<00:00, 2.82kB/s]

vocab.txt:      232k/? [00:00<00:00, 8.41MB/s]
Device set to use cuda:0
[{'label': 'POSITIVE', 'score': 0.999874472618103}]
```

```
from transformers import pipeline
```

```
classifier = pipeline('sentiment-analysis')
result = classifier('The food is not tasty')
print(result)
```



```
No model was supplied, defaulted to distilbert/dis
Using a pipeline without specifying a model name a
Device set to use cuda:0
[{'label': 'NEGATIVE', 'score': 0.9997794032096863}]
```

```
!pip install gradio
```



```
Requirement already satisfied: gradio in /usr/loca
Requirement already satisfied: aiofiles<25.0,>=22.
Requirement already satisfied: anyio<5.0,>=3.0 in
Requirement already satisfied: brotli>=1.1.0 in /u
Requirement already satisfied: fastapi<1.0,>=0.115
Requirement already satisfied: ffmpeg in /usr/local
Requirement already satisfied: gradio-client==1.11
Requirement already satisfied: groovy~=0.1 in /usr
Requirement already satisfied: httpx<1.0,>=0.24.1
Requirement already satisfied: huggingface-hub>=0.
Requirement already satisfied: jinja2<4.0 in /usr/
Requirement already satisfied: markupsafe<4.0,>=2.
Requirement already satisfied: numpy<3.0,>=1.0 in
Requirement already satisfied: orjson~=3.0 in /usr
Requirement already satisfied: packaging in /usr/l
Requirement already satisfied: pandas<3.0,>=1.0 in
Requirement already satisfied: pillow<12.0,>=8.0 i
Requirement already satisfied: pydantic<2.12,>=2.0
Requirement already satisfied: pydub in /usr/local
Requirement already satisfied: python-multipart>=0
Requirement already satisfied: pyyaml<7.0,>=5.0 in
Requirement already satisfied: ruff>=0.9.3 in /usr
Requirement already satisfied: safehttpx<0.2.0,>=0
Requirement already satisfied: semantic-version~=2
Requirement already satisfied: starlette<1.0,>=0.4
Requirement already satisfied: tomlkit<0.14.0,>=0.
Requirement already satisfied: typer<1.0,>=0.12 in
Requirement already satisfied: typing-extensions~=
Requirement already satisfied: uvicorn>=0.14.0 in
Requirement already satisfied: fsspec in /usr/loca
Requirement already satisfied: websockets<16.0,>=1
Requirement already satisfied: idna>=2.8 in /usr/l
Requirement already satisfied: sniffio>=1.1 in /us
Requirement already satisfied: certifi in /usr/loc
Requirement already satisfied: httpcore==1.* in /u
Requirement already satisfied: h11>=0.16 in /usr/l
Requirement already satisfied: filelock in /usr/lo
Requirement already satisfied: requests in /usr/lo
Requirement already satisfied: tqdm>=4.42.1 in /us
Requirement already satisfied: hf-xet<2.0.0,>=1.1.
Requirement already satisfied: python-dateutil>=2.
Requirement already satisfied: pytz>=2020.1 in /us
Requirement already satisfied: tzdata>=2022.7 in /
Requirement already satisfied: annotated-types>=0.
Requirement already satisfied: pydantic-core==2.33
Requirement already satisfied: typing-inspection>=
Requirement already satisfied: click>=8.0.0 in /us
Requirement already satisfied: shellingham>=1.3.0
Requirement already satisfied: rich>=10.11.0 in /u
Requirement already satisfied: six>=1.5 in /usr/lo
Requirement already satisfied: markdown-it-py>=2.2
Requirement already satisfied: pygments<3.0.0,>=2.
Requirement already satisfied: charset-normalizer<
Requirement already satisfied: urllib3<3,>=1.21.1
Requirement already satisfied: mdurl~=0.1 in /usr/
```

```
import gradio as gr
```



```
from transformers import pipeline

classifier = pipeline('sentiment-analysis')

def analyze_sentiment(text):
    return classifier(text)[0]['label']

gr.Interface(fn = analyze_sentiment, inputs='text', outp
```



No model was supplied, defaulted to distilbert/dis
Using a pipeline without specifying a model name a
Device set to use cuda:0
It looks like you are running Gradio on a hosted J

Colab notebook detected. To show errors in colab n
* Running on public URL: <https://acef6eee696505d6a>

This share link expires in 1 week. For free perman

text

The move is very thrilling,i enjoyed it

Clear

Submit

output

POSITIVE

Flag

Use via API · Built with Gradio · Settings

4. Text to Image model using Difussion Pipeline

```
from diffusers import StableDiffusionPipeline
import torch
```

```

model_id = "sd-legacy/stable-diffusion-v1-5"
pipe = StableDiffusionPipeline.from_pretrained(model_id)
pipe = pipe.to("cuda")

prompt = "beautiful girl walking in street light with"
image = pipe(prompt).images[0]

image.save("beautifull girl.png")

```



```

model_index.json: 100%      541/541 [00:00<00:00, 60.3kB/s]

Fetching 15 files: 100%      15/15 [00:30<00:00, 2.22s/it]

preprocessor_config.json: 100%  342/342 [00:00<00:00, 22.8kB/s]

config.json: 100%           617/617 [00:00<00:00, 16.1kB/s]

special_tokens_map.json: 100%  472/472 [00:00<00:00, 15.5kB/s]

config.json:      4.72k/? [00:00<00:00, 117kB/s]

merges.txt:       525k/? [00:00<00:00, 4.14MB/s]

scheduler_config.json: 100%  308/308 [00:00<00:00, 4.35kB/s]

model.safetensors: 100%  1.22G/1.22G [00:19<00:00, 119MB/s]

tokenizer_config.json: 100%   806/806 [00:00<00:00, 82.0kB/s]

model.safetensors: 100%  492M/492M [00:11<00:00, 56.0MB/s]

vocab.json:       1.06M/? [00:00<00:00, 10.4MB/s]

diffusion_pytorch_model.safetensors: 100%  3.44G/3.44G [00:30<00:00, 119MB/s]

config.json: 100%           547/547 [00:00<00:00, 33.1kB/s]

diffusion_pytorch_model.safetensors: 100%  335M/335M [00:07<00:00, 56.0MB/s]

config.json: 100%           743/743 [00:00<00:00, 18.7kB/s]

```

```

import gradio as gr
from diffusers import StableDiffusionPipeline
import torch

# Load model (fp16 + CUDA for performance)
model_id = "sd-legacy/stable-diffusion-v1-5"
pipe = StableDiffusionPipeline.from_pretrained(
    model_id,
    torch_dtype=torch.float16,
    use_auth_token=True # if needed
)
pipe = pipe.to("cuda")

```

```
# Define Gradio interface function
def generate_image(prompt):
    image = pipe(prompt).images[0]
    return image

# UI
gr.Interface(
    fn=generate_image,
    inputs=gr.Textbox(label="Enter your prompt", placeholder="Type a description and generate an image"),
    outputs=gr.Image(type="pil", label="Generated Image"),
    title='Priyanka Text - Image Generator',
    description="Type a description and generate an image using Stable Diffusion v1.5",
).launch()
```



Keyword arguments {'use_auth_token': True} are not supported.
Loading pipeline components...: 100% 7/7 [00:21<00:00, 3.3

It looks like you are running Gradio on a hosted Jupyter Notebook.

Colab notebook detected. To show errors in colab notebook, click the icon in the top bar.

* Running on public URL: <https://19a3e8935ecc5a5f8.gradio.live>

This share link expires in 1 week. For free permanent links, use the upgrade plan.

Priyanka Text - Image Generator

Type a description and generate an image using
Stable Diffusion v1.5

Enter your prompt

e.g. a beautiful girl with blue eyes walking
on the road

Clear

Submit

 Generated Image

Priyanka Text - Image Generator

Type a description and generate an image using Stable Diffusion v1.5

Enter your prompt

Peacock dancing in rain with features

Clear

Submit

Generated Image

